

Are There Fundamental Limitations in Supporting Vector Data Management in Relational Databases?

A Case Study of PostgreSQL

Yunan Zhang · Shige Liu · Jianguo Wang
ICDE 2024

Presenter: Lim Yuan Jee

National Cheng Kung University



Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Introduction

- Vector data is widely used in ML, IR, and data science.
- Comparison between:
 - Specialized Vector Databases (e.g., Faiss, Milvus).
 - Generalized Vector Databases (e.g., PASE, pgvector).
- Motivation: Are there fundamental limitations in relational database to support vector data management.
- Research Question: Can relational databases effectively support vector data management?

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Methodology

- Approach: Compare relational database against specialized vector database
 - PASE (PostgreSQL)
 - FAISS
- Metrics:
 - Index Construction Time
 - Query Time
 - Index Size
 - Recall
- Dataset: Include benchmarks like SIFT1M, GIST1M, and Deep10M

Datasets used

Dataset	# Dimensions	# Vectors	# Queries
SIFT1M [51]	128	1,000,000	10,000
GIST1M [51]	960	1,000,000	1,000
Deep1M [8]	256	1,000,000	1,000
SIFT10M [51]	128	10,000,000	10,000
Deep10M [8]	96	10,000,000	10,000
TURING10M [52]	100	10,000,000	10,000

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- **Key Findings**
 - **Index Construction**
 - Index Size
 - QPS Performance
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Key Findings: Index Construction

- Observations:
 - PASE is 6x-80x slower than Faiss for index construction.
 - Reason: Lack of SGEMM optimization in PASE.




Fig. 3: Index Construction Time for IVF_FLAT




Fig. 4: Index Construction Time for IVF_FLAT Without SGEMM

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- **Key Findings**
 - Index Construction
 - **Index Size**
 - QPS Performance
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Key Findings: Index Size

- Observations:
 - PASE consumes 3x-13x more space for HNSW due to page-based storage.
 - Minimal difference in index size for IVF FLAT and IVF PQ.

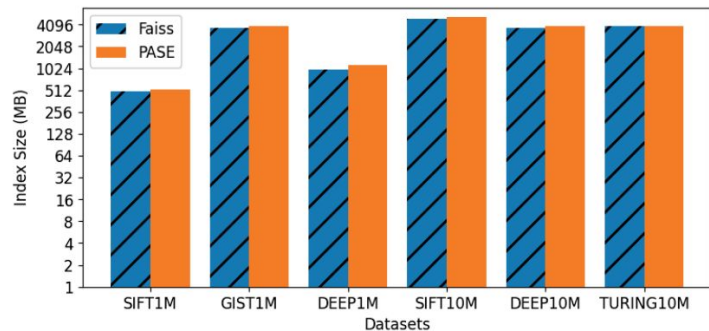


Fig. 11: Index Size for IVF_FLAT




Fig. 13: Index Size for HNSW

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- **Key Findings**
 - Index Construction
 - Index Size
 - **QPS Performance**
- Insights
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Key Findings: QPS Performance

- Observations:
 - PASE query time is 2x-7x slower than Faiss.
 - Performance bottlenecks: Memory management, top-k heap size.




Fig. 14: Search Time for IVF_FLAT




Fig. 17: Search time for HNSW

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- **Insights**
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Insights: Root Causes

- Root causes of performance gap
 - SGEMM Optimization
 - Memory Management
 - Parallel Execution
 - Memory centric Page Structure
 - K-means implementation
 - Heap Size in Top-k Computation.
 - Precomputed Table Implementation




Fig. 15: Search Time for IVF_FLAT With Replaced Centroids (Faiss*)

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- **Insights**
 - Root Causes
 - **Impact of SGEMM**
 - Impact of parallelism
- Conclusion

Insights: Impact of SGEMM (Single Precision General Matrix Multiplication)

- SGEMM plays a crucial role in optimizing computations within vector databases, particularly for distance calculations and K-means clustering.
- Why is it important?
 - Matrix Operations in Vector Databases
- Challenges Without SGEMM:
 - Computations are performed using basic loops or less efficient matrix operations.
 - ex: Faiss uses SGEMM to transform the distance calculation problem into a matrix-matrix multiplication

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- **Insights**
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

▶ Insights: Impact of Parallelism

- Index Construction
 - PASE lacks efficient multi-threading during the index construction phase
- Query Execution
 - PASE Query Bottleneck
 - Searching multiple buckets simultaneously or updating shared data structures (e.g., heaps) become bottlenecks.
 - Faiss utilizes local heaps for parallel searching within buckets.
- Results
 - Faiss demonstrates near-linear scalability with thread count during query execution

Discussion topics

- Introduction
- Methodology
 - Dataset Used
- Key Findings
 - Index Construction
 - Index Size
 - QPS Performance
- **Insights**
 - Root Causes
 - Impact of SGEMM
 - Impact of parallelism
- Conclusion

Conclusion

- Conclusion: No fundamental limitations; performance gaps are due to implementation.
- Future Work:
 - Build a memory-centric generalized vector database.
 - Improve algorithmic and parallelism support.
 - High priority for SGEMM
 - Parallelism
 - Optimized top k computation